

NAME: PRATYUSH ABHI

CAMPUS: V.I.T VELLORE

We are building the project on Mnist Dataset

About Mnist data: The MNIST database contains 60,000 training images and 10,000 testing images. Half of the training set and half of the test set were taken from NIST's training dataset, while the other half of the training set and the other half of the test set were taken from NIST's testing dataset.

Here we would be using ANN and CNN models to train and test the data sets we are getting an test accuracy of 97.73%

```
In [ ]: import torch import torch.nn
as nn import
torch.nn.functional as
```

```
In [2]: from torch.utils.data import DataLoader
from torchvision import
datasets, transforms
```

```
In [3]: import numpy as np
import pandas as pd
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [4]: #MNIST image to tensor
```

```
In [5]: transform=transforms.ToTensor()
```

```
In [6]: train_data=datasets.MNIST(root='/Users/pratyushabhi/Desktop/3rd Year/
```

```
In [7]: test_data=datasets.MNIST(root='/Users/pratyushabhi/Desktop/3rd Year/I
```

```
In [8]: train_data
```

```
Out[8]: Dataset MNIST
      Number of datapoints: 60000
      Root location: /Users/pratyushabhi/Desktop/3rd Year/Image process
ing/ANN
      Split: Train
      StandardTransform
      Transform: ToTensor()
```

```
In [9]: test_data
```

```
Out[9]: Dataset MNIST
```



```
0.9922, 0.9922, 0.8824, 0.6745, 0.9922, 0.9490, 0.7647,
0.2510 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.1922 ,
0.9333, 0.9922, 0.9922, 0.9922, 0.9922, 0.9922, 0.9922,
0.9922 ,
0.9922, 0.9843, 0.3647, 0.3216, 0.3216, 0.2196, 0.1529,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0706 ,
0.8588, 0.9922, 0.9922, 0.9922, 0.9922, 0.9922, 0.7765,
0.7137 ,
0.9686, 0.9451, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.3137, 0.6118, 0.4196, 0.9922, 0.9922, 0.8039, 0.0431,
0.0000 ,
0.1686, 0.6039, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0549, 0.0039, 0.6039, 0.9922, 0.3529, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.5451, 0.9922, 0.7451, 0.0078,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0431, 0.7451, 0.9922, 0.2745,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.1373, 0.9451, 0.8824,
0.6275 ,
0.4235, 0.0039, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.3176, 0.9412,
0.9922 ,
0.9922, 0.4667, 0.0980, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.1765,
0.7294 ,
```

```
0.9922, 0.9922, 0.5882, 0.1059, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0627 ,
0.3647, 0.9882, 0.9922, 0.7333, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.9765, 0.9922, 0.9765, 0.2510, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.1804,
0.5098 ,
0.7176, 0.9922, 0.9922, 0.8118, 0.0078, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.1529, 0.5804, 0.8980,
0.9922 ,
0.9922, 0.9922, 0.9804, 0.7137, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0941, 0.4471, 0.8667, 0.9922, 0.9922,
0.9922 ,
0.9922, 0.7882, 0.3059, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0902, 0.2588, 0.8353, 0.9922, 0.9922, 0.9922, 0.9922,
0.7765 ,
0.3176, 0.0078, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0706,
0.6706 ,
0.8588, 0.9922, 0.9922, 0.9922, 0.9922, 0.7647, 0.3137,
0.0353 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.2157, 0.6745, 0.8863, 0.9922 ,
0.9922, 0.9922, 0.9922, 0.9569, 0.5216, 0.0431, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.5333, 0.9922, 0.9922, 0.9922 ,
0.8314, 0.5294, 0.5176, 0.0627, 0.0000, 0.0000, 0.0000,
0.0000 ,
```

```

0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000], [0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000 ,
0.0000, 0.0000, 0.0000, 0.0000]])),
5)

```

```
In [12]: image,label=train_data[0]
```

```
In [13]: image.shape
```


```
Out[13]: torch.Size([1, 28, 28])
```

```
In [14]: label
```

```
Out[14]: 5
```


```
In [15]: plt.imshow(image.reshape((28,28)))
```

```
Out[15]: <matplotlib.image.AxesImage at 0x163a81ed0>
```



```
In [16]: plt.imshow(image.reshape((28,28)), cmap='gray')
```

```
Out[16]: <matplotlib.image.AxesImage at 0x163ae6710>
```



```

In [17]: torch.manual_seed(101)

train_loader=DataLoader(train_data,batch_size=100,shuffle=True)
test_loader=DataLoader(test_data,batch_size=500,shuffle=False)

In [18]: from torchvision.utils import make_grid
np.set_printoptions(formatter=dict(int=lambda x:f'{x:4}')) #FORMATTING

In [19]: #First Batch for images,labels in
train_loader:
    break

In [20]: images.shape # 1->Means it is a grey scale image,(28,28)->width,length o

Out[20]: torch.Size([100, 1, 28, 28])

In [21]: labels.shape


Out[21]:

#Print first 12 labels
print('Labels:',labels[:12].numpy())

#Print the first 12 images
im=make_grid(images[:12],nrow=12)#default nrow is 8
plt.figure(figsize=(10,4))
plt.imshow(np.transpose(im.numpy(),(1,2,0)));

torch.Size([100])
In [22]:

```



```
In [23]: class MultilayerPerceptron(nn.Module):
    def
    __init__(self, in_sz=784, out_sz=10, layers=[120, 84]):
        super().__init__()

        self.fc1=nn.Linear(in_sz, layers[0])
    self.fc2=nn.Linear(layers[0], layers[1])
    self.fc3=nn.Linear(layers[1], out_sz)

    def forward(self, x): #x->data
        features          x=F.relu(self.fc1(x))
    x=F.relu(self.fc2(x))
    x=self.fc3(x)

    return F.log_softmax(x, dim=1) #Multi-class classification
```

```
In [24]: torch.manual_seed(101)
    model=MultilayerPerceptron()
    model
```

```
Out[24]: MultilayerPerceptron(
  (fc1): Linear(in_features=784, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84,
    bias=True)  (fc3): Linear(in_features=84, out_features=10,
    bias=True) )
```

```
In [25]: criterion=nn.CrossEntropyLoss()
    optimizer=torch.optim.Adam(model.parameters(), lr=0.001)
```

```
In [26]: images.shape
```

```
Out[26]: torch.Size([100, 1, 28, 28])
```

```
In [27]: images.view(100, -1)
```

```
Out[27]: tensor([[0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 ...,
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.]])
```



```

In [28]: import time
start_time = time.time()

epochs = 10
train_losses = []
test_losses = []
train_correct = []
test_correct = []

for i in range(epochs):
    trn_corr = 0
    tst_corr = 0

    # Run the training batches
    for b, (X_train,
y_train) in enumerate(train_loader):
        b+=1

        # Apply the model
        y_pred = model(X_train.view(100, -1)) # Here we flatten X_train
        loss = criterion(y_pred, y_train)

        # Tally the number of correct predictions
        predicted = torch.max(y_pred.data, 1)[1]
        batch_corr = (predicted == y_train).sum()
        trn_corr += batch_corr

        # Update parameters
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Print interim results
        if b%200 == 0:
            print(f'epoch: {i:2}  batch: {b:4} [{100*b:6}/60000]  loss:
{ accuracy: {trn_corr.item()*100/(100*b):7.3f}%}')

        # Update train loss & accuracy for the epoch
        train_losses.append(loss.item())
        train_correct.append(trn_corr.item())

    # Run the testing batches
    with torch.no_grad():
        for b, (X_test, y_test) in enumerate(test_loader):

            # Apply the model
            y_val = model(X_test.view(500, -1)) # Here we flatten X_test

            # Tally the number of correct predictions
            predicted = torch.max(y_val.data, 1)[1]

```

```

        tst_corr += (predicted == y_test).sum()

        # Update test loss & accuracy for the epoch
        loss = criterion(y_val, y_test)
        test_losses.append(loss)
        test_correct.append(tst_corr)

    print(f'\nDuration: {time.time() - start_time:.0f} seconds') # print the

#total_time=time.time()-start_time
#print(f'Duration:{Total time/60}mins')

```

```

epoch:  0  batch:  200 [ 20000/60000]  loss: 0.23562382  accuracy:
83.24
5%
epoch:  0  batch:  400 [ 40000/60000]  loss: 0.35330707  accuracy:
87.53
2%
epoch:  0  batch:  600 [ 60000/60000]  loss: 0.13765770  accuracy:
89.58
2%
epoch:  1  batch:  200 [ 20000/60000]  loss: 0.24507998  accuracy:
94.88
0%
epoch:  1  batch:  400 [ 40000/60000]  loss: 0.14064841  accuracy:
95.11
0%
epoch:  1  batch:  600 [ 60000/60000]  loss: 0.07430533  accuracy:
95.37
7%
epoch:  2  batch:  200 [ 20000/60000]  loss: 0.08479684  accuracy:
96.57
0%
epoch:  2  batch:  400 [ 40000/60000]  loss: 0.06338982  accuracy:
96.67
2%
epoch:  2  batch:  600 [ 60000/60000]  loss: 0.06284785  accuracy:
96.73
7%
epoch:  3  batch:  200 [ 20000/60000]  loss: 0.11593810  accuracy:
97.65
0%
epoch:  3  batch:  400 [ 40000/60000]  loss: 0.05100821  accuracy:
97.46
8%
epoch:  3  batch:  600 [ 60000/60000]  loss: 0.03086828  accuracy:
97.49
3%
epoch:  4  batch:  200 [ 20000/60000]  loss: 0.11303577  accuracy:
98.17
0%

```

```

epoch: 4 batch: 400 [ 40000/60000] loss: 0.04967898 accuracy:
98.01
3%
epoch: 4 batch: 600 [ 60000/60000] loss: 0.05145194 accuracy:
98.01
0%
epoch: 5 batch: 200 [ 20000/60000] loss: 0.00721604 accuracy:
98.41
0%
epoch: 5 batch: 400 [ 40000/60000] loss: 0.03383062 accuracy:
98.46
8%
epoch: 5 batch: 600 [ 60000/60000] loss: 0.05648751 accuracy:
98.38
0%
epoch: 6 batch: 200 [ 20000/60000] loss: 0.02997145 accuracy:
98.82
0%
epoch: 6 batch: 400 [ 40000/60000] loss: 0.06433750 accuracy:
98.69
0%
epoch: 6 batch: 600 [ 60000/60000] loss: 0.08998419 accuracy:
98.67 0%
epoch: 7 batch: 200 [ 20000/60000] loss: 0.07536934 accuracy:
99.10
5%
epoch: 7 batch: 400 [ 40000/60000] loss: 0.09858016 accuracy:
98.91
0%
epoch: 7 batch: 600 [ 60000/60000] loss: 0.03994036 accuracy:
98.86
5%
epoch: 8 batch: 200 [ 20000/60000] loss: 0.00415698 accuracy:
99.22
0%
epoch: 8 batch: 400 [ 40000/60000] loss: 0.00953338 accuracy:
99.13
5%
epoch: 8 batch: 600 [ 60000/60000] loss: 0.08832055 accuracy:
99.11
0%
epoch: 9 batch: 200 [ 20000/60000] loss: 0.00536544 accuracy:
99.34
0%
epoch: 9 batch: 400 [ 40000/60000] loss: 0.01200775 accuracy:
99.26
5%
epoch: 9 batch: 600 [ 60000/60000] loss: 0.02104353 accuracy:
99.24 2%

```

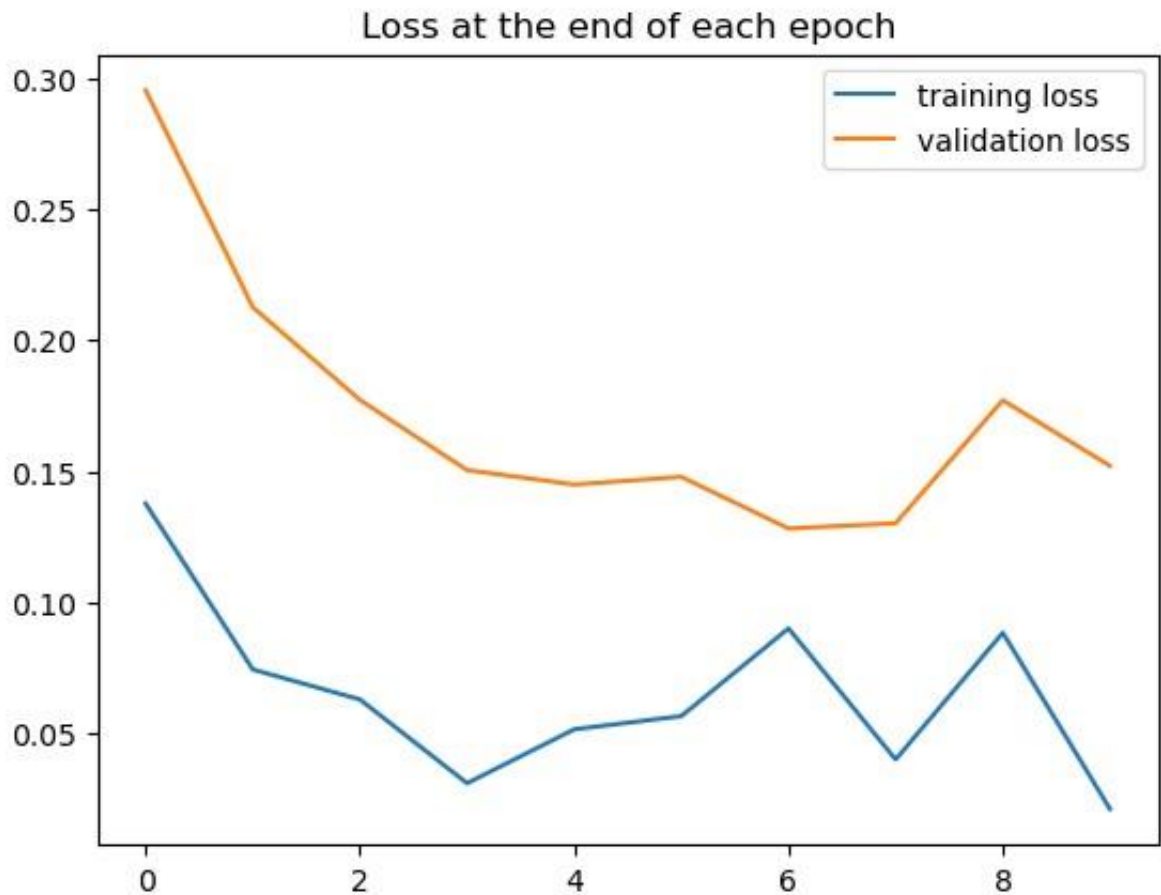
Duration: 39 seconds

```

In [29]: plt.plot(train_losses, label='training loss')
         plt.plot(test_losses, label='validation loss')

```

```
plt.title('Loss at the end of each epoch')
plt.legend();
```



```
print(test_correct) # contains the results of all 10 epochs
print()
print(f'Test accuracy: {test_correct[-1].item()*100/10000:.3f}%') # print
```

Evaluating test data

In [30]:

```
[tensor(9437), tensor(9579), tensor(9691), tensor(9706), tensor(9746),
 tensor(9759), tensor(9777), tensor(9756), tensor(9751), tensor(9773)]
```

Test accuracy: 97.730%

```
In [31]: # Extract the data all at once, not in batches test_load_all =
         DataLoader(test_data, batch_size=10000, shuffle=False)
```

```
In [32]: with torch.no_grad(): correct = 0 for X_test, y_test in
         test_load_all: y_val = model(X_test.view(len(X_test), -1)) #
         pass in a flattene predicted = torch.max(y_val, 1)[1]
         correct += (predicted == y_test).sum() print(f'Test accuracy:
         {correct.item()}/{len(test_data)} = {correct.item(
```

Test accuracy: 9773/10000 = 97.730% Confusion

Matrix

In [33]: *# print a row of values for reference*

```
np.set_printoptions(formatter=dict(int=lambda x: f'{x:4}'))
print(np.arange(10).reshape(1,10)) print()
```

print the confusion matrix

```
print(confusion_matrix(predicted.view(-1), y_test.view(-1)))
```

```
[[ 0  1  2  3  4  5  6  7  8  9]]
```

```
[[ 970  0  4  0  1  2  6  3  5  3]
 [  0 1124  4  0  0  0  3  7  0  2]
 [  1  3 1002  2  1  0  2 11  2  0]
 [  3  3  5 991  0 10  1  3  4  2]
 [  1  0  2  0 961  3  4  2  2  5]
 [  0  0  0  5  0 865  3  1  2  2]
 [  1  1  2  0  6  2 938  0  1  0]
 [  1  0  3  4  1  0  0 978  3  0]
 [  2  4 10  2  1  8  1  2 952  3]
 [  1  0  0  6 11  2  0 21  3 992]]
```

```
import torch import torch.nn as nn import
torch.nn.functional as F from
torch.utils.data import DataLoader from
torchvision import datasets, transforms
from torchvision.utils import make_grid

import numpy as np
import pandas as pd
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
%matplotlib inline
```

Using CNN for "Mnist" Dataset

In [34]:

```
In [36]: transform = transforms.ToTensor()
```

```
train_data = datasets.MNIST(root='/Users/pratyushabhi/Desktop/3rd Yea  
test_data = datasets.MNIST(root='/Users/pratyushabhi/Desktop/3rd Year
```

```
In [37]: train_data
```

```
Out[37]: Dataset MNIST  
        Number of datapoints: 60000  
        Root location: /Users/pratyushabhi/Desktop/3rd Year/Image process  
ing/ANN  
        Split: Train  
        StandardTransform  
        Transform: ToTensor()
```

```
In [38]: test_data
```

```
Out[38]: Dataset MNIST  
        Number of datapoints: 10000  
        Root location: /Users/pratyushabhi/Desktop/3rd Year/Image process  
ing/ANN      Split: Test  
        StandardTransform  
        Transform: ToTensor()
```

```
In [39]: train_loader = DataLoader(train_data, batch_size=10, shuffle=True)  
test_loader = DataLoader(test_data, batch_size=10, shuffle=False)
```

```
In [40]:
```

```
#Defining CNN
```

```
In [41]:
```

```
In [42]: conv1 = nn.Conv2d(1, 6, 3, 1)  
conv2 = nn.Conv2d(6, 16, 3, 1)
```

```
In [43]: # Grab the first MNIST record for i, (X_train,  
y_train) in enumerate(train_data):  
        break
```

```
# Create a rank-4 tensor to be passed into the model  
# (train_loader will have done this  
already) x = X_train.view(1,1,28,28)  
print(x.shape)  
torch.Size([1, 1, 28, 28])
```

```
In [44]: # Perform the first  
convolution/activation x =  
F.relu(conv1(x)) print(x.shape)  
torch.Size([1, 6, 26, 26])
```

```
In [45]: # Run the first pooling layer
x = F.max_pool2d(x, 2, 2)
print(x.shape)

torch.Size([1, 6, 13, 13])
```

```
In [46]: # Perform the second
convolution/activation x =
F.relu(conv2(x)) print(x.shape)

torch.Size([1, 16, 11, 11])
```

```
In [48]: # Run the second pooling layer
x = F.max_pool2d(x, 2, 2)
print(x.shape)

torch.Size([1, 16, 5, 5])
```

```
In [49]: # Flatten the data x
= x.view(-1, 5*5*16)
print(x.shape)

torch.Size([1, 400])
```

```
In [50]: class ConvolutionalNetwork(nn.Module):
def __init__(self):
super().__init__()
self.conv1 = nn.Conv2d(1, 6, 3, 1)
self.conv2 = nn.Conv2d(6, 16, 3, 1)
self.fc1 = nn.Linear(5*5*16, 120)
self.fc2 = nn.Linear(120, 84)
self.fc3 = nn.Linear(84, 10)

def forward(self, X):
X = F.relu(self.conv1(X))
X = F.max_pool2d(X, 2, 2)
X = F.relu(self.conv2(X))
X = F.max_pool2d(X, 2, 2)
X = X.view(-1, 5*5*16)
X = F.relu(self.fc1(X))
X = F.relu(self.fc2(X))
X = self.fc3(X)
return F.log_softmax(X, dim=1)
```

```
In [51]: torch.manual_seed(42) model =
ConvolutionalNetwork() model
```

```
Out[51]: ConvolutionalNetwork(
  (conv1): Conv2d(1, 6, kernel_size=(3, 3), stride=(1, 1))
  (conv2): Conv2d(6, 16, kernel_size=(3, 3), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)
```

```
In [52]: #Loss function
```

```
In [53]: criterion = nn.CrossEntropyLoss()  
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
```



```

In [54]: #Train the model
import time
start_time = time.time()

epochs = 5
train_losses = []
test_losses = []
train_correct = []
test_correct = []

for i in range(epochs):
    trn_corr = 0
    tst_corr = 0

    # Run the training batches    for b, (X_train,
y_train) in enumerate(train_loader):
        b+=1

        # Apply the model
        y_pred = model(X_train) # we don't flatten X-train here
        loss = criterion(y_pred, y_train)

        # Tally the number of correct predictions
        predicted = torch.max(y_pred.data, 1)[1]
        batch_corr = (predicted == y_train).sum()
        trn_corr += batch_corr

        # Update parameters
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Print interim results
        if b%600 == 0:
            print(f'epoch: {i:2}  batch: {b:4} [{10*b:6}/60000]  loss:
{1 accuracy: {trn_corr.item()*100/(10*b):7.3f}%')

            train_losses.append(loss.item())
            train_correct.append(trn_corr.item())

        # Run the testing batches    with torch.no_grad():
        for b, (X_test, y_test) in enumerate(test_loader):

            # Apply the model
            y_val = model(X_test)

            # Tally the number of correct predictions
            predicted = torch.max(y_val.data, 1)[1]
            tst_corr += (predicted == y_test).sum()

            loss = criterion(y_val, y_test)
            test_losses.append(loss)
            test_correct.append(tst_corr)

        print(f'\nDuration: {time.time() - start_time:.0f} seconds') # print the

```

```
epoch: 0 batch: 600 [ 6000/60000] loss: 0.04055630 accuracy:
78.41
7%
epoch: 0 batch: 1200 [ 12000/60000] loss: 0.08253471 accuracy:
85.80
0%
epoch: 0 batch: 1800 [ 18000/60000] loss: 0.36470532 accuracy:
88.68
9%
epoch: 0 batch: 2400 [ 24000/60000] loss: 0.01825019 accuracy:
90.52
5%
epoch: 0 batch: 3000 [ 30000/60000] loss: 0.00806712 accuracy:
91.65
0%
epoch: 0 batch: 3600 [ 36000/60000] loss: 0.00097706 accuracy:
92.49
2%
epoch: 0 batch: 4200 [ 42000/60000] loss: 0.44326892 accuracy:
93.13
3%
epoch: 0 batch: 4800 [ 48000/60000] loss: 0.03169333 accuracy:
93.61
5%
epoch: 0 batch: 5400 [ 54000/60000] loss: 0.01946524 accuracy:
94.03
1%
epoch: 0 batch: 6000 [ 60000/60000] loss: 0.02709320 accuracy:
94.33
3%
epoch: 1 batch: 600 [ 6000/60000] loss: 0.01472266 accuracy:
97.75
0%
epoch: 1 batch: 1200 [ 12000/60000] loss: 0.04359249 accuracy:
97.87
5%
epoch: 1 batch: 1800 [ 18000/60000] loss: 0.00124075 accuracy:
97.90
0%
epoch: 1 batch: 2400 [ 24000/60000] loss: 0.03912879 accuracy:
97.85
4%
epoch: 1 batch: 3000 [ 30000/60000] loss: 0.14564939 accuracy:
97.86
7%
epoch: 1 batch: 3600 [ 36000/60000] loss: 0.00049980 accuracy:
97.87
8%
epoch: 1 batch: 4200 [ 42000/60000] loss: 0.00076085 accuracy:
97.91
4%
epoch: 1 batch: 4800 [ 48000/60000] loss: 0.00105086 accuracy:
97.91
2%
```

```
epoch: 1 batch: 5400 [ 54000/60000] loss: 0.00745581 accuracy:
97.94
1%
epoch: 1 batch: 6000 [ 60000/60000] loss: 0.13721663 accuracy:
97.94
3%
epoch: 2 batch: 600 [ 6000/60000] loss: 0.00099742 accuracy:
98.68
3%
epoch: 2 batch: 1200 [ 12000/60000] loss: 0.00254112 accuracy:
98.65
0%
epoch: 2 batch: 1800 [ 18000/60000] loss: 0.00188525 accuracy:
98.53
9%
epoch: 2 batch: 2400 [ 24000/60000] loss: 0.00276521 accuracy:
98.59
6%
epoch: 2 batch: 3000 [ 30000/60000] loss: 0.24948892 accuracy:
98.55
3%
epoch: 2 batch: 3600 [ 36000/60000] loss: 0.03415573 accuracy:
98.54
4%
epoch: 2 batch: 4200 [ 42000/60000] loss: 0.02346098 accuracy:
98.52
4%
epoch: 2 batch: 4800 [ 48000/60000] loss: 0.01159347 accuracy:
98.54
6%
epoch: 2 batch: 5400 [ 54000/60000] loss: 0.00048137 accuracy:
98.53 1%
epoch: 2 batch: 6000 [ 60000/60000] loss: 0.00013920 accuracy:
98.53
5%
epoch: 3 batch: 600 [ 6000/60000] loss: 0.00057079 accuracy:
98.90
0%
epoch: 3 batch: 1200 [ 12000/60000] loss: 0.00067071 accuracy:
98.80
0%
epoch: 3 batch: 1800 [ 18000/60000] loss: 0.00059610 accuracy:
98.86
1%
epoch: 3 batch: 2400 [ 24000/60000] loss: 0.00048164 accuracy:
98.81
7%
epoch: 3 batch: 3000 [ 30000/60000] loss: 0.12320199 accuracy:
98.83
3%
epoch: 3 batch: 3600 [ 36000/60000] loss: 0.00446792 accuracy:
98.83
3%
epoch: 3 batch: 4200 [ 42000/60000] loss: 0.00079797 accuracy:
98.84
```

```

8%
epoch: 3 batch: 4800 [ 48000/60000] loss: 0.12382223 accuracy:
98.84
2%
epoch: 3 batch: 5400 [ 54000/60000] loss: 0.00634904 accuracy:
98.83
3%
epoch: 3 batch: 6000 [ 60000/60000] loss: 0.00734646 accuracy:
98.85
3%
epoch: 4 batch: 600 [ 6000/60000] loss: 0.00278111 accuracy:
99.08
3%
epoch: 4 batch: 1200 [ 12000/60000] loss: 0.17682204 accuracy:
99.06
7%
epoch: 4 batch: 1800 [ 18000/60000] loss: 0.00004966 accuracy:
99.08
3%
epoch: 4 batch: 2400 [ 24000/60000] loss: 0.00025236 accuracy:
99.06
7%
epoch: 4 batch: 3000 [ 30000/60000] loss: 0.00490724 accuracy:
99.04
0%
epoch: 4 batch: 3600 [ 36000/60000] loss: 0.05752410 accuracy:
99.04
4%
epoch: 4 batch: 4200 [ 42000/60000] loss: 0.01335440 accuracy:
99.07
1%
epoch: 4 batch: 4800 [ 48000/60000] loss: 0.00020917 accuracy:
99.06
5%
epoch: 4 batch: 5400 [ 54000/60000] loss: 0.00006763 accuracy:
99.05
0%
epoch: 4 batch: 6000 [ 60000/60000] loss: 0.00072345 accuracy:
99.05 7%

```


Duration: 74 seconds

In [55]: *## Plot the loss and accuracy comparisons*

```

In [56]: plt.plot(train_losses, label='training loss')
plt.plot(test_losses, label='validation loss')
plt.title('Loss at the end of each epoch')
plt.legend();


```



```
In [57]: test_losses
```

```
Out[57]: [tensor(0.0023),  
          tensor(0.0024),  
          tensor(0.0008),  
          tensor(0.0001),  
          tensor(6.7299e-05)]
```

```
In [58]: plt.plot([t/600 for t in train_correct], label='training accuracy')  
         plt.plot([t/100 for t in test_correct], label='validation  
         accuracy') plt.title('Accuracy at the end of each epoch')  
         plt.legend();
```



```
In [59]: #Evaluate Test data
test_load_all = DataLoader(test_data, batch_size=10000, shuffle=False)
with torch.no_grad():
    correct = 0
    for X_test, y_test in test_load_all:
        y_val = model(X_test) # we don't flatten the data this time
        predicted = torch.max(y_val,1)[1]
        correct += (predicted == y_test).sum()
print(f'Test accuracy: {correct.item()}/{len(test_data)} = {correct.item() / len(test_data):.4f}')
```

Test accuracy: 9862/10000 = 98.620%

```
In [64]: #Confusion Matrix
# print a row of values for reference
np.set_printoptions(formatter=dict(int=lambda x: f'{x:4}'))
print(np.arange(10).reshape(1,10))

# print the confusion matrix
print(confusion_matrix(predicted.view(-1), y_test.view(-1)))
```

```
[[ 0  1  2  3  4  5  6  7  8  9]]


[[ 975  0  0  0  0  3  4  0  4  0]
 [  0 1134  4  1  1  0  4  5  0  3]
 [  2  0 1019  1  0  0  0  5  2  0]
 [  0  1  0 1002  0 19  0  0  1  0]
 [  0  0  1  0 978  0  2  0  1 16]
 [  0  0  0  2  0 858  0  0  0  2]]
```

```
[ 1  0  0  0  2  2 947  0  0  1]
[ 1  0  6  2  0  0  0 1015  3  6]
[ 1  0  2  2  0  3  1  1 962  9]
[ 0  0  0  0  1  7  0  2  1 972]]
```

Run a new image through the model.

We can also pass a single image through the model to obtain a prediction. Pick a number from 0 to 9999, assign it to "x", and we'll use that value to select a number from the MNIST test set.

```
In [61]: x = 2019
plt.figure(figsize=(1,1))
plt.imshow(test_data[x][0].reshape((28,28)), cmap="gist_yarg");
```



```
In [62]: model.eval() with
torch.no_grad():
    new_pred = model(test_data[x][0].view(1,1,28,28)).argmax()
    print("Predicted value:", new_pred.item())
```

Predicted value: 9